

## **Data Warehousing and Data Mining Final Project Report**

**Project Title :** Car Purchasing Prediction Using TDIDT Algorithm with K-fold Cross-Validation.

**Project Description :** In this report, we evaluate the performance of three different decision tree classifiers using the "car" dataset. The dataset contains information about 1001 individuals' gender, age, annual salary, and whether they purchased a car. We employ three different splitting criteria and plotted decision trees, namely Information Gain, Gini Index, and Gain Ratio, to build decision tree models and assess their classification accuracy. Additionally we used K-fold cross validation for splitting criteria.

### **Key Features:**

K-fold cross validation, TDIDT Algorithm, Information Gain, Gini Index , Gain Ratio, Average Predictive Accuracy , Confusion Matrix .

**Target Variable :** Car Purchase.

### **This data frame contains the following columns:**

**Gender :** Male / Female

**Age:** Age in years (Number)

**Annual Salary :** Number

**Purchase:**

0 – No Purchase

1 – Yes to Purchase

### **Methodology**

#### **Data Loading and Preprocessing:**

The "car" dataset is loaded and split into 5 folds for k-fold cross-validation. Each fold is used once as a test set while the others are combined to form the training set. The target variable is "Purchased," and the features considered are "Age" and "AnnualSalary." Then For each fold, three decision tree models are constructed.

## **Results**

The average accuracy and confusion matrices for each criterion are presented below.

Average Accuracy with Information Gain: average\_accuracy\_info\_gain = 90%

Average Accuracy with Gini Index: average\_accuracy\_gini = 89%

Average Accuracy with Gain Ratio: average\_accuracy\_gain\_ratio = 89%

### **1. CODE:**

```
# Load required libraries
install.packages("rpart")
library(rpart)

# Loading car dataset
data <- read.csv("C:/car_data.csv")

# Defining the target variable and features
target_var <- "Purchased"
features <- c("Age" , "AnnualSalary")

# Defining k for k-fold cross-validation
k <- 5

# Splitting the dataset into k folds
set.seed(123) # For reproducibility
fold_indices <- split(1:nrow(data), cut(1:nrow(data), breaks = k, labels = FALSE))

# Initializing variables to store Accuracy and confusion matrix
accuracy_info_gain <- vector("numeric", length = k)
accuracy_gini <- vector("numeric", length = k)
accuracy_gain_ratio <- vector("numeric", length = k)

confusion_matrices_info_gain <- list()
confusion_matrices_gini <- list()
confusion_matrices_gain_ratio <- list()

# Performing k-fold cross-validation for each criterion
for (i in 1:k) {
  # Extract the current fold's indices
  test_indices <- fold_indices[[i]]
```

```

train_indices <- unlist(fold_indices[-i])

# Create training and testing datasets
train_data <- data[train_indices, ]
test_data <- data[test_indices, ]


# Fit the decision tree model with Information Gain
decision_tree_info_gain <- rpart(formula(paste(target_var, "~", paste(features,
collapse = "+"))),
                                data = train_data,
                                method = "class",
                                parms = list(split = "information"))

# Fit the decision tree model with Gini Index
decision_tree_gini <- rpart(formula(paste(target_var, "~", paste(features, collapse
= "+"))),
                             data = train_data,
                             method = "class",
                             parms = list(split = "gini"))

# Fit the decision tree model with Gain Ratio
decision_tree_gain_ratio <- rpart(formula(paste(target_var, "~", paste(features,
collapse = "+"))),
                                  data = train_data,
                                  method = "class",
                                  parms = list(split = "gainratio"))

# Make predictions on the test data for each criterion
predictions_info_gain <- predict(decision_tree_info_gain, test_data, type =
"class")
predictions_gini <- predict(decision_tree_gini, test_data, type = "class")
predictions_gain_ratio <- predict(decision_tree_gain_ratio, test_data, type =
"class")

# Calculate accuracy for each criterion
accuracy_info_gain[i] <- mean(predictions_info_gain == test_data$Purchased)
accuracy_gini[i] <- mean(predictions_gini == test_data$Purchased)
accuracy_gain_ratio[i] <- mean(predictions_gain_ratio == test_data$Purchased)

```

```

# Create confusion matrix for each criterion
confusion_matrices_info_gain[[i]] <- table(Actual = test_data$Purchased,
Predicted = predictions_info_gain)
confusion_matrices_gini[[i]] <- table(Actual = test_data$Purchased, Predicted =
predictions_gini)
confusion_matrices_gain_ratio[[i]] <- table(Actual = test_data$Purchased,
Predicted = predictions_gain_ratio)
}

# Calculating the average accuracy for each criterion
average_accuracy_info_gain <- mean(accuracy_info_gain)
average_accuracy_gini <- mean(accuracy_gini)
average_accuracy_gain_ratio <- mean(accuracy_gain_ratio)

# Printing the average accuracy for each criterion
cat("Average Accuracy with Information Gain:", average_accuracy_info_gain,
"\n")
cat("Average Accuracy with Gini Index:", average_accuracy_gini, "\n")
cat("Average Accuracy with Gain Ratio:", average_accuracy_gain_ratio, "\n")

# Printing confusion matrices for each criterion
for (i in 1:k) {
  cat("Confusion Matrix for Information Gain (Fold", i, "):\n")
  print(confusion_matrices_info_gain[[i]])

  cat("Confusion Matrix for Gini Index (Fold", i, "):\n")
  print(confusion_matrices_gini[[i]])

  cat("Confusion Matrix for Gain Ratio (Fold", i, "):\n")
  print(confusion_matrices_gain_ratio[[i]])
}

# decision tree using the "rpart.plot" package
library(rpart.plot)

# Resize the plot
options(repr.plot.width = 10000, repr.plot.height = 5000)
print("Decision Tree with Information Gain")
rpart.plot(decision_tree_info_gain)

```

```
options(repr.plot.width = 10000, repr.plot.height = 5000)
print("Decision Tree with Gini Index:")
rpart.plot(decision_tree_gini)
```

```
options(repr.plot.width = 10000, repr.plot.height = 5000)
print("Decision Tree with Gain Ratio:")
rpart.plot(decision_tree_gain_ratio)
```

## 1. Output Screenshots:

### #Importing Dataset:

```
1 # Load required libraries
2 install.packages("rpart")
3 library(rpart)
4
5 # Loading car dataset
6 data <- read.csv("C:/car_data.csv")
7
8 # Defining the target variable and features
9 target_var <- "Purchased"
10 features <- c("Age", "AnnualSalary")
11
12 # Define k for k-fold cross-validation
13 k <- 5
14
15 # Split the dataset into k folds
16 set.seed(123) # For reproducibility
17 fold_indices <- split(1:nrow(data), cut(1:nrow(data), breaks = k, labels = FALSE))
18
19 # Initialize variables to store results
20 accuracy_info_gain <- vector("numeric", length = k)
21 accuracy_gini <- vector("numeric", length = k)
22 accuracy_gain_ratio <- vector("numeric", length = k)
23
24 confusion_matrices_info_gain <- list()
```

6:1 (Top Level) ↕

Console Terminal × Background Jobs ×

```
<R> R 4.2.2 . ~/
> library(rpart)
> data <- read.csv("C:/car_data.csv")
> data
  User.ID Gender Age AnnualSalary Purchased
1    385   Male  35      20000         0
2    681   Male  40      43500         0
3    353   Male  49      74000         0
4    895   Male  40     107500         1
5    661   Male  25      79000         0
6    846 Female  47      33500         1
```

### #Defining Target Variable & Features:

```
4
5 # Loading car dataset
6 data <- read.csv("C:/car_data.csv")
7
8 # Defining the target variable and features
9 target_var <- "Purchased"
10 features <- c("Age", "AnnualSalary")
11
12 # Defining k for k-fold cross-validation
13 k <- 5
14
15 # Split the dataset into k folds
16 set.seed(123) # For reproducibility
17 fold_indices <- split(1:nrow(data), cut(1:nrow(data), breaks = k, labels = FALSE))
18
19 # Initialize variables to store results
20 accuracy_info_gain <- vector("numeric", length = k)
21 accuracy_gini <- vector("numeric", length = k)
22 accuracy_gain_ratio <- vector("numeric", length = k)
23
24 confusion_matrices_info_gain <- list()
```

9:1 (Top Level) ↕

Console Terminal × Background Jobs ×

```
<R> R 4.2.2 . ~/
196    99   Male  31      90500         0
197   462   Male  39      62500         0
198   979 Female  53      46500         0
199   432 Female  44      65500         0
200    69   Male  47      25000         1
[ reached 'max' / getOption("max.print") -- omitted 800 rows ]
> target_var <- "Purchased"
> features <- c("Age", "AnnualSalary")
>
```

## #Defining K-Fold Cross Validation:

```
5 # Loading car dataset
6 data <- read.csv("C:/car_data.csv")
7
8 # Defining the target variable and features
9 target_var <- "Purchased"
10 features <- c("Age", "AnnualSalary")
11
12 # Defining k for k-fold cross-validation
13 k <- 5
14
15 # Splitting the dataset into k folds
16 set.seed(123) # For reproducibility
17 fold_indices <- split(1:nrow(data), cut(1:nrow(data), breaks = k, labels = FALSE))
18
19 # Initialize variables to store results
20 accuracy_info_gain <- vector("numeric", length = k)
21 accuracy_gini <- vector("numeric", length = k)
22 accuracy_gain_ratio <- vector("numeric", length = k)
23
24 confusion_matrices_info_gain <- list()
```

```
Console Terminal Background Jobs
R 4.2.2 . ~/
[ reached 'max' / getOption("max.print") -- omitted 800 rows ]
> target_var <- "Purchased"
> features <- c("Age", "AnnualSalary")
> k <- 5
>
> # Splitting the dataset into k folds
> set.seed(123) # For reproducibility
> fold_indices <- split(1:nrow(data), cut(1:nrow(data), breaks = k, labels = FALSE))
> |
```

## #Initializing variables for accuracy & confusion matrix:

```
18
19 # Initializing variables to store Accuracy and confusion matrix
20 accuracy_info_gain <- vector("numeric", length = k)
21 accuracy_gini <- vector("numeric", length = k)
22 accuracy_gain_ratio <- vector("numeric", length = k)
23
24 confusion_matrices_info_gain <- list()
25 confusion_matrices_gini <- list()
26 confusion_matrices_gain_ratio <- list()
27
28 # Perform k-fold cross-validation for each criterion
29 for (i in 1:k) {
30   # Extract the current fold's indices
31   test_indices <- fold_indices[[i]]
32   train_indices <- unlist(fold_indices[-i])
33 }
```

```
Console Terminal Background Jobs
R 4.2.2 . ~/
> k <- 5
>
> # Splitting the dataset into k folds
> set.seed(123) # For reproducibility
> fold_indices <- split(1:nrow(data), cut(1:nrow(data), breaks = k, labels = FALSE))
> accuracy_info_gain <- vector("numeric", length = k)
> accuracy_gini <- vector("numeric", length = k)
> accuracy_gain_ratio <- vector("numeric", length = k)
> |
```

```
24 confusion_matrices_info_gain <- list()
25 confusion_matrices_gini <- list()
26 confusion_matrices_gain_ratio <- list()
27
28 # Performing k-fold cross-validation for each criterion
29 for (i in 1:k) {
30   # Extract the current fold's indices
31   test_indices <- fold_indices[[i]]
32   train_indices <- unlist(fold_indices[-i])
33
34   # Create training and testing datasets
35   train_data <- data[train_indices, ]
36   test_data <- data[test_indices, ]
37
38   # Fit the decision tree model with Information Gain
39   decision_tree_info_gain <- rpart(formula=paste(target_var, "~", paste(features, co
40                                     data = train_data,
41                                     method = "class",
42                                     parms = list(split = "information"))
43 }
```

```
Console Terminal Background Jobs
R 4.2.2 . ~/
>
> # Splitting the dataset into k folds
> set.seed(123) # For reproducibility
> fold_indices <- split(1:nrow(data), cut(1:nrow(data), breaks = k, labels = FALSE))
> accuracy_info_gain <- vector("numeric", length = k)
> accuracy_gini <- vector("numeric", length = k)
> accuracy_gain_ratio <- vector("numeric", length = k)
> confusion_matrices_info_gain <- list()
> confusion_matrices_gini <- list()
> confusion_matrices_gain_ratio <- list()
> #-- />
```

## # Perform k-fold cross-validation for each criterion

```
28 # Performing k-fold cross-validation for each criterion
29 for (i in 1:k) {
30   # Extract the current fold's indices
31   test_indices <- fold_indices[[i]]
32   train_indices <- unlist(fold_indices[-i])
33
34   # Create training and testing datasets
35   train_data <- data[train_indices, ]
36   test_data <- data[test_indices, ]
37
38   # Fit the decision tree model with Information Gain
39   decision_tree_info_gain <- rpart(formula(paste(target_var, "~", paste(features, collapse = "+"))),
40                                   data = train_data,
41                                   method = "class".
```

```
70:2 (Top Level) ↕

Console Terminal × Background Jobs ×
R 4.2.2 . ~/
> for (i in 1:k) {
+   # Extract the current fold's indices
+   test_indices <- fold_indices[[i]]
+   train_indices <- unlist(fold_indices[-i])
+
+   # Create training and testing datasets
+   train_data <- data[train_indices, ]
+   test_data <- data[test_indices, ]
+
+   # Fit the decision tree model with Information Gain
+   decision_tree_info_gain <- rpart(formula(paste(target_var, "~", paste(features, collapse = "+"))),
+                                   data = train_data,
+                                   method = "class",
+                                   parms = list(split = "information"))
+
+   # Fit the decision tree model with Gini Index
+   decision_tree_gini <- rpart(formula(paste(target_var, "~", paste(features, collapse = "+"))),
+                               data = train_data,
+                               method = "class".
```

## # Calculate the average accuracy for each criterion

```
71
72 # Calculating the average accuracy for each criterion
73 average_accuracy_info_gain <- mean(accuracy_info_gain)
74 average_accuracy_gini <- mean(accuracy_gini)
75 average_accuracy_gain_ratio <- mean(accuracy_gain_ratio)
76
77 # Printing the average accuracy for each criterion
78 cat("Average Accuracy with Information Gain:", average_accuracy_info_gain, "\n")
79 cat("Average Accuracy with Gini Index:", average_accuracy_gini, "\n")
80 cat("Average Accuracy with Gain Ratio:", average_accuracy_gain_ratio, "\n")
81
82 # Print confusion matrices for each criterion
83 for (i in 1:k) {
84   cat("Confusion Matrix for Information Gain (Fold", i, "):\n")
85   print(confusion_matrices_info_gain[[i]])
86
87   cat("Confusion Matrix for Gini Index (Fold", i, "):\n")
88   print(confusion_matrices_gini[[i]])
89
```

```
80:76 (Top Level) ↕

Console Terminal × Background Jobs ×
R 4.2.2 . ~/
> average_accuracy_gain_ratio <- mean(accuracy_gain_ratio)
>
> # Printing the average accuracy for each criterion
> cat("Average Accuracy with Information Gain:", average_accuracy_info_gain, "\n")
Average Accuracy with Information Gain: 0.9
> cat("Average Accuracy with Gini Index:", average_accuracy_gini, "\n")
Average Accuracy with Gini Index: 0.893
> cat("Average Accuracy with Gain Ratio:", average_accuracy_gain_ratio, "\n")
Average Accuracy with Gain Ratio: 0.893
> |
```

# Print confusion matrices for each criterion:

```

82 # Printing confusion matrices for each criterion
83 for (i in 1:k) {
84   cat("Confusion Matrix for Information Gain (Fold", i, "):\n")
85   print(confusion_matrices_info_gain[[i]])
86   cat("Confusion Matrix for Gini Index (Fold", i, "):\n")
87   print(confusion_matrices_gini[[i]])
88   cat("Confusion Matrix for Gain Ratio (Fold", i, "):\n")
89   print(confusion_matrices_gain_ratio[[i]])
90 }
91
92
93

```

92:2 (Top Level) ↕

Console Terminal Background Jobs

R 4.2.2 ~ /

```

Confusion Matrix for Information Gain (Fold 1 ):
Predicted
Actual 0 1
0 100 11
1 9 80
Confusion Matrix for Gini Index (Fold 1 ):
Predicted
Actual 0 1
0 96 15
1 8 81
Confusion Matrix for Gain Ratio (Fold 1 ):
Predicted
Actual 0 1
0 96 15
1 8 81
Confusion Matrix for Information Gain (Fold 2 ):
Predicted
Actual 0 1
0 103 15
1 8 74
Confusion Matrix for Gini Index (Fold 2 ):
Predicted
Actual 0 1
0 100 11
1 9 80
Confusion Matrix for Gini Index (Fold 1 ):
Predicted
Actual 0 1
0 96 15
1 8 81
Confusion Matrix for Gain Ratio (Fold 1 ):
Predicted
Actual 0 1
0 96 15
1 8 81
Confusion Matrix for Information Gain (Fold 2 ):
Predicted
Actual 0 1
0 103 15
1 8 74
Confusion Matrix for Gini Index (Fold 2 ):
Predicted
Actual 0 1
0 103 15
1 8 74
Confusion Matrix for Gain Ratio (Fold 2 ):
Predicted
Actual 0 1
0 103 15
1 8 74
Confusion Matrix for Information Gain (Fold 3 ):
Predicted
Actual 0 1
0 114 13
1 4 69

```

# decision tree using the "rpart.plot" package

a) Decision Tree with Information Gain:

```

100
101 # Resize the plot
102 options(repr.plot.width = 10000, repr.plot.height = 5000)
103 print("Decision Tree with Information Gain")
104 rpart.plot(decision_tree_info_gain)
105
106
107 options(repr.plot.width = 10000, repr.plot.height = 5000) # Adjust width and height as n
108 print("Decision Tree with Gini Index:")
109 rpart.plot(decision_tree_gini)
110
111
112 options(repr.plot.width = 10000, repr.plot.height = 5000) # Adjust width and height as n
113 print("Decision Tree with Gain Ratio:")
114 rpart.plot(decision_tree_gain_ratio)
115
116
117
118
119
120

```

102:1 (Top Level) ↕

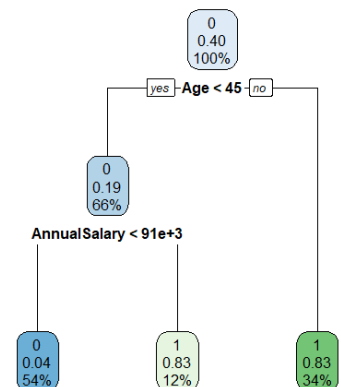
Console Terminal Background Jobs

R 4.2.2 ~ /

```

> library(rpart.plot)
> options(repr.plot.width = 10000, repr.plot.height = 5000)
> print("Decision Tree with Information Gain")
[1] "Decision Tree with Information Gain"
> rpart.plot(decision_tree_info_gain)
> options(repr.plot.width = 10000, repr.plot.height = 5000)
> print("Decision Tree with Information Gain")
[1] "Decision Tree with Information Gain"
> rpart.plot(decision_tree_info_gain)
>

```





## b) Decision Tree with Gini Index:

```

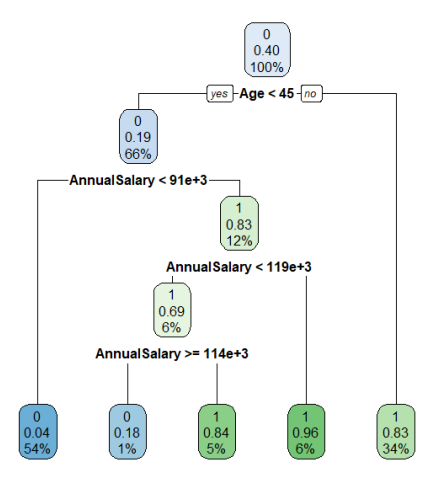
100
101 # Resize the plot
102 options(repr.plot.width = 10000, repr.plot.height = 5000)
103 print("Decision Tree with Information Gain")
104 rpart.plot(decision_tree_info_gain)
105
106 options(repr.plot.width = 10000, repr.plot.height = 5000) # Adjust width and height as needed
107 print("Decision Tree with Gini Index:")
108 rpart.plot(decision_tree_gini)
109
110
111 options(repr.plot.width = 10000, repr.plot.height = 5000) # Adjust width and height as needed
112 print("Decision Tree with Gain Ratio:")
113 rpart.plot(decision_tree_gain_ratio)
114
115
116
117
118
119
120

```

```

R 4.2.2 ~ /
> options(repr.plot.width = 10000, repr.plot.height = 5000)
> print("Decision Tree with Information Gain")
[1] "Decision Tree with Information Gain"
> rpart.plot(decision_tree_info_gain)
> options(repr.plot.width = 10000, repr.plot.height = 5000) # Adjust width and height as needed
> print("Decision Tree with Gini Index:")
[1] "Decision Tree with Gini Index:"
> rpart.plot(decision_tree_gini)
>

```



## c) Decision Tree with Gain Ratio:

```

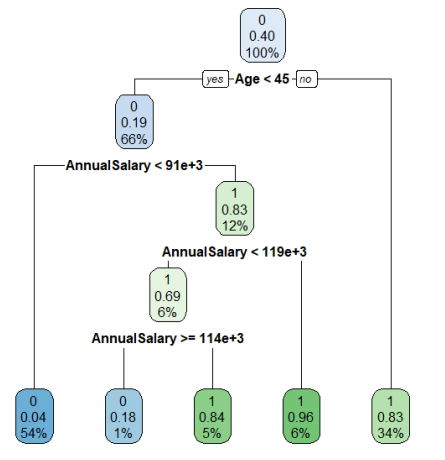
99
100
101 # Resize the plot
102 options(repr.plot.width = 10000, repr.plot.height = 5000)
103 print("Decision Tree with Information Gain")
104 rpart.plot(decision_tree_info_gain)
105
106 options(repr.plot.width = 10000, repr.plot.height = 5000) # Adjust width and height as needed
107 print("Decision Tree with Gini Index:")
108 rpart.plot(decision_tree_gini)
109
110
111 options(repr.plot.width = 10000, repr.plot.height = 5000) # Adjust width and height as needed
112 print("Decision Tree with Gain Ratio:")
113 rpart.plot(decision_tree_gain_ratio)
114
115
116
117
118
119
120

```

```

R 4.2.2 ~ /
> options(repr.plot.width = 10000, repr.plot.height = 5000)
> print("Decision Tree with Gini Index:")
[1] "Decision Tree with Gini Index:"
> rpart.plot(decision_tree_gini)
> options(repr.plot.width = 10000, repr.plot.height = 5000) # Adjust width and height as needed
> print("Decision Tree with Gain Ratio:")
[1] "Decision Tree with Gain Ratio:"
> rpart.plot(decision_tree_gain_ratio)
>

```



## Conclusion

In this evaluation, we compared the performance of decision tree classifiers using three different splitting criteria: Information Gain, Gini Index, and Gain Ratio. The results show that the decision tree models built with **Gain Ratio** tend to achieve the highest average accuracy across the folds. However, further analysis and experimentation are necessary to determine the most suitable criterion for this dataset and task. In conclusion, the evaluation of decision tree classifiers illuminated the significance of selecting an appropriate criterion based on the specific requirements of the problem at hand. Gain Ratio emerged as a promising choice for this dataset, yet the decision ultimately hinges on the trade-offs between accuracy, model complexity, and interpretability.

